

# Biostat276 Project 2

Zhaodong Wu

2/16/2022

## Contents

<b>Bayesian Probit Regression</b>	<b>1</b>
Frequentist Method (Not related to the questions, just a try)	2
Question 1	2
Question 2	5
Question 3	9

## Bayesian Probit Regression

In R load the package `(survival)` and consider the analysis of the data-set `(infert)`. Ignoring dependence due to matching, consider a Bayesian analysis for a logistic regression model relating case status to: age, parity, education, spontaneous and induced. More precisely, assume case status  $y_i$  has density  $y_i \sim_{ind} Bern(p_i)$ ,  $p_i = \Phi(X_i'\beta)$ , where  $\Phi(\cdot)$  is the standard Gaussian cdf. Consider a prior  $\beta \sim N(0, 10^2(X'X)^{-1})$ . We are interested in  $p(\beta|Y)$ .

```
rm(list = ls())
library(ggplot2)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.0 --

## v tibble 3.1.0      v dplyr 1.0.7
## v tidyr 1.1.3       v stringr 1.4.0
## v readr 1.4.0       v forcats 0.5.1
## v purrr 0.3.4

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(mvtnorm)
library(kableExtra)

##
## Attaching package: 'kableExtra'

## The following object is masked from 'package:dplyr':
##
##     group_rows

library(truncnorm)

# load the data
library(survival)
```

```
data("infert")

# change the data type of `education`
inferthw <- inferth
inferthw$education <- as.numeric(inferthw$education) - 1
dat <- inferthw %>%
  dplyr::select(c("education", "age", "parity", "induced", "case",
                  "spontaneous"))
str(dat) # all variables are numeric

## 'data.frame': 248 obs. of 6 variables:
## $ education : num 0 0 0 0 1 1 1 1 1 1 ...
## $ age : num 26 42 39 34 35 36 23 32 21 28 ...
## $ parity : num 6 1 6 4 3 4 1 2 1 2 ...
## $ induced : num 1 1 2 2 1 2 0 0 0 0 ...
## $ case : num 1 1 1 1 1 1 1 1 1 1 ...
## $ spontaneous: num 2 0 0 0 1 1 0 0 1 0 ...
```

## Frequentist Method (Not related to the questions, just a try)

Before answering the questions, we can directly compute the MLE of the model parameters using `glm()` function with a probit link function.

```
freq <- glm(case ~ age + parity + education + spontaneous + induced,
            data = dat, family = binomial(link = probit))

# MLE estimator
freq$coefficients

## (Intercept)      age      parity  education spontaneous      induced
## -0.91127035  0.02073028 -0.43967798 -0.28754507  1.16488031  0.71911970

# Deviance
freq$deviance # the deviance is large

## [1] 259.6669
```

## Question 1

(1) Describe and implement an adaptive Metropolis-Hastings algorithm designed to obtain a MC with stationary distribution  $p(\beta|Y)$ .

The full posterior distribution for the Bayesian binary probit model can be computed as follows:

$$\begin{aligned}
 \pi(\beta|Y, X) &\propto \pi(\beta) \cdot \pi(Y, X|\beta) \\
 &= \pi(\beta) \cdot \prod_{i=1}^n p(y_i, X_i|\beta) \\
 &= \pi(\beta) \cdot \prod_{i=1}^n \Phi(X'_i \beta)^{y_i} [1 - \Phi(X'_i \beta)]^{1-y_i} \\
 &\propto \exp\left[-\frac{\beta'(X'X)\beta}{200}\right] \cdot \prod_{i=1}^n \Phi(X'_i \beta)^{y_i} [1 - \Phi(X'_i \beta)]^{1-y_i}
 \end{aligned}$$

It's obvious that  $\pi(\beta)$  is not a conjugate prior by the fact that no conjugate prior  $\pi(\beta)$  exists for the parameters of the probit regression model.

We compute the posterior distribution first:

```
# Calculation of posterior
beta_prior <- 100 * solve(as.matrix(t(dat[, -5])) %*% as.matrix(dat[, -5]))

posterior <- function(beta) {
  pi <- pnorm(as.matrix(dat[, -5]) %*% t(beta))
  data <- data.frame(pi = pi, y = dat[, 5])
  post <- apply(data, 1, function(x) {ifelse(x[2] == 1, x[1], 1 - x[1])})
  post <- log(post) %>%
    sum(.) - 1/2 * beta %*% solve(beta_prior) %*% t(beta)
}

# Adaptive Random Walk-Metropolis Hastings
mh.Q1 <- function(nsim = 10000, burn = 0.2, # chain parameters
                  delta = 0.75, c = 1,      # set c = 1 (c > 0)
                  seed = 1998) {
  # initialization-----
  set.seed(seed)
  nsim1 <- nsim * (1 + burn)
  burni <- nsim * burn

  beta_num <- dim(dat)[2] - 1

  beta <- matrix(data = rep(0, beta_num), nrow = 1)
  beta.ch <- matrix(data = NA, nrow = nsim, ncol = beta_num)
  betavar <- diag(beta_num) * (10^(-14)) # error term \epsilon
  deltaset <- rbinom(nsim1, 1, delta)
  # run chain-----

  for (i in 1:nsim1) {
    # here we use adaptive MH for the first 2200 iterations
    if (i <= 2200) {
      # \Sigma_{t}^{\tilde{}} = \Sigma_t + \epsilon * I
      betavar <- (betavar * (i - 1) + t(beta) %*% beta)/i +
        diag(beta_num) * (10^(-14))
      delta_tm <- deltaset[i] # delta = 0.75 during adaptive
    } else {delta_tm <- 1} # delta is fixed after adaptive
    # use the last variance of beta from adaptive for the remaining iterations
    beta_tm1 <- rmvnorm(n = 1, mean = beta, sigma = c * betavar)
    beta_tm2 <- rmvnorm(n = 1, mean = beta, sigma = beta_prior)
    beta_tm <- beta_tm1 * delta_tm + beta_tm2 * (1 - delta_tm)

    P0 <- posterior(beta)
    P1 <- posterior(beta_tm)
    ratio <- P1 - P0
    if (log(runif(1)) < ratio) {
      beta <- beta_tm
    }
    # Store Chain after burn
    if (i > burni) {
      i1 <- i - burni
    }
  }
}
```


```

    beta.ch[i1, ] <- beta
  }
}
return(list(beta = beta.ch))
}

# Simulation
arw_mh <- mh.Q1(nsim = 10000, burn = 0.2, delta = 0.75, c = 1)

# convergence checking
par(mfrow = c(2, 3))
for (i in 1:5) {
  plot(1:10000, arw_mh$beta[, i], type = "l",
       ylab = substitute(beta[x], list(x = i)))
}

```



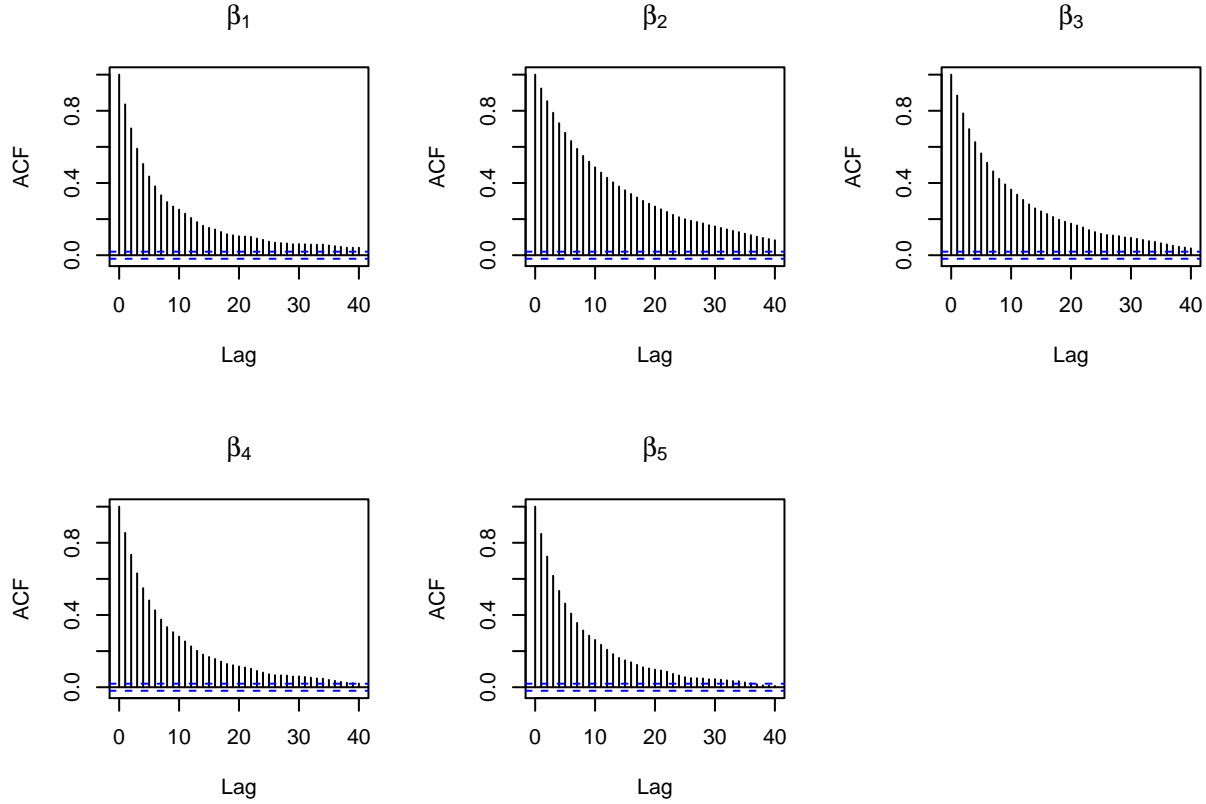
```

# Autocorrelation plot
par(mfrow = c(2, 3))
for (i in 1:5) {
  acf(arw_mh$beta[, i], main = substitute(beta[x], list(x = i)))
} # autocorrelation plots look good (beta_2 is kind of worse but acceptable)

```

Table 1: Summary Table of Coefficients Using Adaptive MH

	education	age	parity	induced	spontaneous
2.5%	-0.760	0.005	-0.662	0.480	0.751
50%	-0.613	0.014	-0.484	0.667	1.033
97.5%	-0.457	0.022	-0.319	0.850	1.319



```
# get the result
result_Q1 <- apply(arw_mh$beta, 2,
  function(x) {quantile(x, c(0.025, 0.5, 0.975))}) %>%
  round(., 3) # 95% credible interval & median
colnames(result_Q1) <- colnames(dat)[-5]

result_Q1 %>%
  kbl(caption = "Summary Table of Coefficients Using Adaptive MH") %>%
  kable_classic(full_width = F, html_font = "Cambria")
```

## Question 2

Describe and implement a data augmented (DA-MCMC) strategy targeting  $p(\beta|y)$ .

Here targeting  $p(\beta, z|y)$  could be easier enough.

Let prior of  $\beta : \beta \sim \mathcal{N}(0, \Sigma_\beta)$ , i.e.  $\Sigma_\beta = 10^2(X'X)^{-1}$

Let  $Z_i|\beta \sim \mathcal{N}(X_i'\beta, 1)$  and define the sampling model conditionally as  $Y_i|Z_i = I(Z_i > 0)$ . Then we use the Gibbs Sampling:

For  $\beta$  :

$$\begin{aligned}
p(\beta|Z_{1:n}, Y_{1:n}) &= p(\beta|Z_{1:n}) \\
&\propto \prod_{i=1}^n \exp\left[-\frac{(z_i - X'_i\beta)^2}{2}\right] \cdot \exp\left(-\frac{\beta^T \Sigma_\beta^{-1} \beta}{2}\right) \\
&= \exp\left[-\frac{(Z - X\beta)'(Z - X\beta) + \beta' \Sigma_\beta^{-1} \beta}{2}\right] \\
&\propto \exp\left[\frac{\beta'(X'X + \Sigma_\beta^{-1})\beta - 2\beta'X'Z}{2}\right]
\end{aligned}$$

By completing the square, we realize that the density is proportional to a normal kernel, the posterior of  $\beta$  satisfies a normal distribution:

$$p(\beta|Z_{1:n}, Y_{1:n}) \sim \mathcal{N}\{(X'X + \Sigma_\beta^{-1})^{-1}X'Z, (X'X + \Sigma_\beta^{-1})^{-1}\}$$

For  $Z_i$  :

$$\begin{aligned}
p(Z_i|Y_i, \beta) &\propto p(Y_i|Z_i) \cdot p(Z_i|\beta) \\
&= I(Z_i > 0) \cdot \exp\left[-\frac{(z_i - X'_i\beta)^2}{2}\right] \quad (\text{if } y_i = 1) \\
&= I(Z_i \leq 0) \cdot \exp\left[-\frac{(z_i - X'_i\beta)^2}{2}\right] \quad (\text{if } y_i = 0)
\end{aligned}$$

The posterior of  $Z_i$  follows a **truncated normal distributon**, i.e.

$$p(Z_i|Y_i, \beta) = \begin{cases} \mathcal{TN}(X'_i\beta, 1, 0, +\infty) & \text{if } y_i = 1 \\ \mathcal{TN}(X'_i\beta, 1, -\infty, 0) & \text{if } y_i = 0 \end{cases}$$

```

DA_Q2 <- function(nsim = 10000, burn = 0.2, # chain parameters
                  seed = 1998,
                  x = as.matrix(dat[-5]), # load the dataset
                  y = as.matrix(dat[5])) {
  # initialization-----
  set.seed(seed)
  nsim1 <- nsim * (1 + burn)
  burn1 <- nsim * burn

  beta_num <- dim(dat)[2] - 1

  beta <- matrix(data = rep(0, beta_num), nrow = 1)
  beta.ch <- matrix(data = NA, nrow = nsim, ncol = beta_num)

  # generate data z-----
  z <- rep(0, length(y))
  z.ch <- matrix(data = NA, nrow = nsim, ncol = length(y))

  range <- cbind(ifelse(y == 1, 0, -Inf),
                 ifelse(y == 1, Inf, 0)
                 )

```

```


# run chain-----
for (i in 1:nsim1) {
  # z
  z_comb <- cbind(x %*% t(beta),
                  1,
                  range)
  z <- apply(z_comb, 1, function(comb) {
    rtruncnorm(1, comb[1], sd = comb[2], a = comb[3], b = comb[4])
  })
  # beta
  beta_mean <- solve(t(x) %*% x + solve(beta_prior)) %*% t(x) %*% z
  beta_var <- solve(t(x) %*% x + solve(beta_prior))
  beta <- rmvnorm(1, mean = beta_mean, sigma = beta_var)

  if (i > burni) {
    i1 <- i - burni
    beta.ch[i1, ] <- beta
  }
}
return(list(beta = beta.ch))
}

# Simulation
DA_Gibbs <- DA_Q2(nsim = 10000, burn = 0.2)

par(mfrow = c(2, 3))
for (i in 1:5) {
  plot(1:10000, DA_Gibbs$beta[, i], type = "l",
       ylab = substitute(beta[x], list(x = i)))
}

```

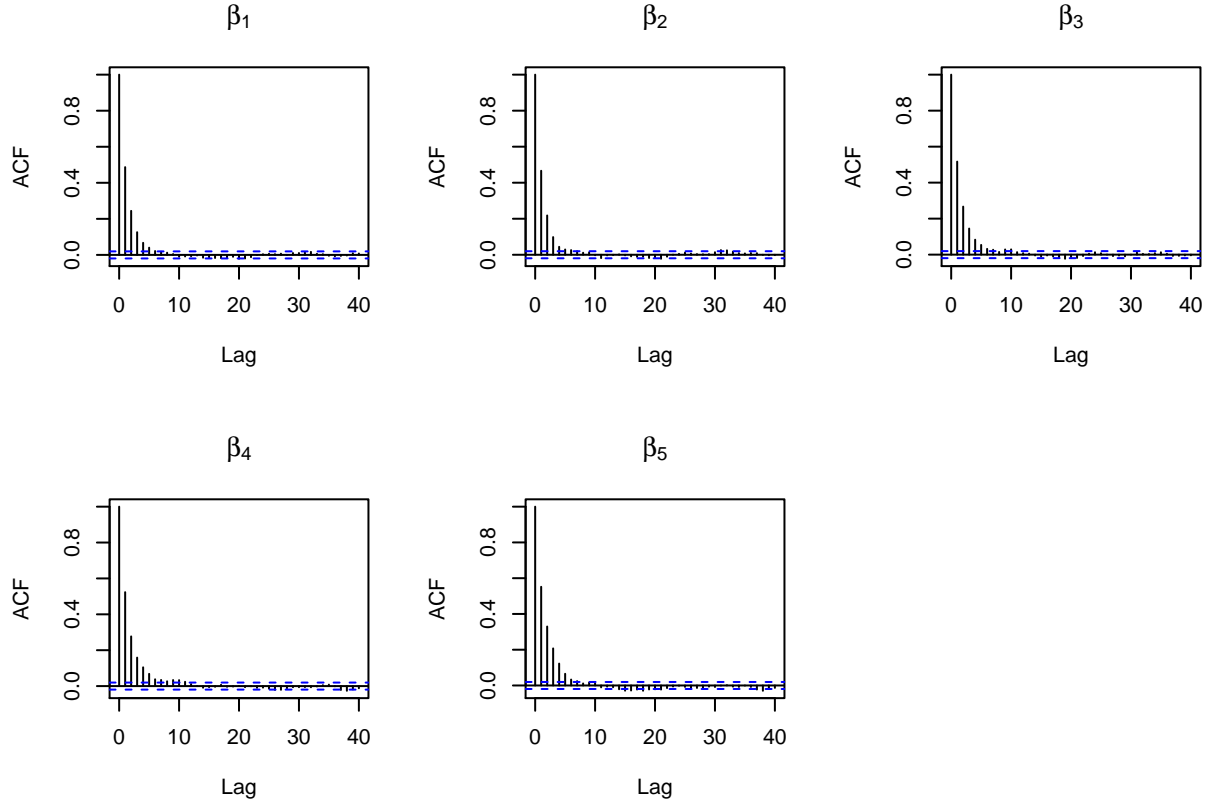


```
# Autocorrelation plot
par(mfrow = c(2, 3))
for (i in 1:5) {
  acf(DA_Gibbs$beta[, i], main = substitute(beta[x], list(x = i)))
} # autocorrelation plots look good (beta_2 is kind of worse but acceptable)
```



Table 2: Summary Table of Coefficients using Data Augmentation

	education	age	parity	induced	spontaneous
2.5%	-0.678	-0.016	-0.656	0.372	0.829
50%	-0.416	0.000	-0.456	0.695	1.141
97.5%	-0.148	0.016	-0.260	1.022	1.463



```
# get the result
result_Q2 <- apply(DA_Gibbs$beta, 2,
  function(x) {quantile(x, c(0.025, 0.5, 0.975))}) %>%
  round(., 3) # 95% credible interval & median
colnames(result_Q2) <- colnames(dat)[-5]

result_Q2 %>%
  kbl(caption = "Summary Table of Coefficients using Data Augmentation") %>%
  kable_classic(full_width = F, html_font = "Cambria")
```

### Question 3

Describe and implement a parameter expanded - data augmentation (PX-DA MCMC) algorithm targeting  $p(\beta|Y)$ .